


Decision Tree and Boosting

David Ivan

June 18, 2018


Decision Tree




Train a decision tree

1000


Train a decision tree




Train a decision tree




Train a decision tree



Train a decision tree




Train a decision tree




How to split a node?

```
max_split_quality = 0
for f in features:
    x_sorted = sort(x[:,f])
    for x_0, x_1 in zip(x_sorted[:-1], x_sorted[1:]):
        midpoint = 0.5*(x_0 + x_1)
        split_quality = get_split_quality(x, f, midpoint)
        if split_quality > max_split_quality:
            max_split_quality = split_quality
            feature_split = f
            threshold = midpoint
return feature_split, threshold
```


Split




Split



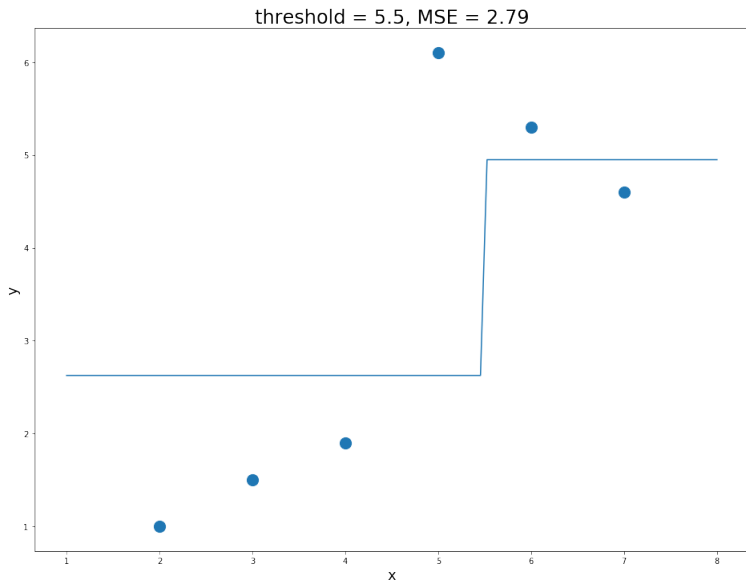
example




example

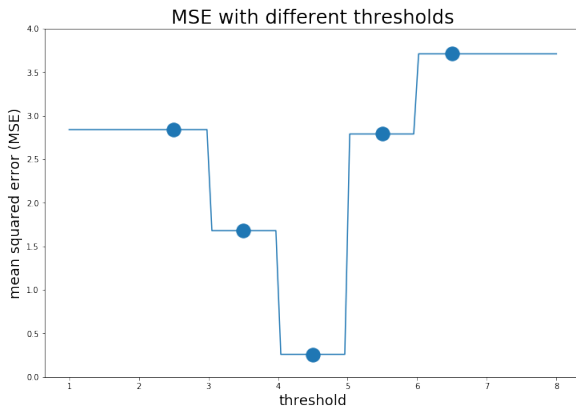



example




example









example



example





example





greedy splits vs optimal splits


greedy split is not optimal in long-term




partitioning the feature space




Classification with decision trees




Classification with decision trees



Classification with decision trees




Classification with decision trees




Impurity measure: entropy

$$H(X) = H(p_1, p_2, \dots, p_n) = -\sum_i^n p_i \cdot \log_2(p_i)$$


Information Theory




$$H(X)=2.58$$



$$H(Y)=2.45$$



$$H(Z)=1.87$$



$$H(V)=0.057$$

Conditional entropy:

$$H(X|Y = y) = -\sum_i p(x_i|y) \cdot \log_2(p(x_i|y))$$

$$H(X|Y) = \sum_y P(Y = y) \cdot H(X|Y = y)$$


Information gain (information conveyed about X by Y):

$$I(X|Y) = H(X) - H(X|Y)$$

Information Theory

$$X = \begin{cases} \text{"red"} & (9/30) \\ \text{"white"} & (8/30) \\ \text{"green"} & (13/30) \end{cases}$$

$$Y = \begin{cases} \text{"left"} & (10/30) \\ \text{"right"} & (20/30) \end{cases}$$




$$(X|Y=\text{"left"}) = \begin{cases} \text{"red"} & (2/10) \\ \text{"white"} & (1/10) \\ \text{"green"} & (7/10) \end{cases}$$

$$(X|Y=\text{"right"}) = \begin{cases} \text{"red"} & (7/20) \\ \text{"white"} & (7/20) \\ \text{"green"} & (6/20) \end{cases}$$

Information Theory

$$X = \begin{cases} \text{"red"} & (9/30) \\ \text{"white"} & (8/30) \\ \text{"green"} & (13/30) \end{cases}$$

$$Y = \begin{cases} \text{"left"} & (10/30) \\ \text{"right"} & (20/30) \end{cases}$$



$$(X|Y=\text{"left"}) = \begin{cases} \text{"red"} & (2/10) \\ \text{"white"} & (1/10) \\ \text{"green"} & (7/10) \end{cases}$$

$$(X|Y=\text{"right"}) = \begin{cases} \text{"red"} & (7/20) \\ \text{"white"} & (7/20) \\ \text{"green"} & (6/20) \end{cases}$$


$$H(X) = H\left(\frac{9}{30}, \frac{8}{30}, \frac{13}{30}\right)$$

$$= -\frac{9}{30} \log_2\left(\frac{9}{30}\right) - \frac{8}{30} \log_2\left(\frac{8}{30}\right) - \frac{13}{30} \log_2\left(\frac{13}{30}\right) = 1.5524$$

Information Theory

$$X = \begin{cases} \text{"red"} & (9/30) \\ \text{"white"} & (8/30) \\ \text{"green"} & (13/30) \end{cases}$$

$$Y = \begin{cases} \text{"left"} & (10/30) \\ \text{"right"} & (20/30) \end{cases}$$



$$(X|Y=\text{"left"}) = \begin{cases} \text{"red"} & (2/10) \\ \text{"white"} & (1/10) \\ \text{"green"} & (7/10) \end{cases}$$

$$(X|Y=\text{"right"}) = \begin{cases} \text{"red"} & (7/20) \\ \text{"white"} & (7/20) \\ \text{"green"} & (6/20) \end{cases}$$


$$H(X|Y = \text{left}) = H\left(\frac{2}{10}, \frac{1}{10}, \frac{7}{10}\right)$$

$$= -\frac{2}{10} \log_2\left(\frac{2}{10}\right) - \frac{1}{10} \log_2\left(\frac{1}{10}\right) - \frac{7}{10} \log_2\left(\frac{7}{10}\right) = 1.1568$$

Information Theory

$$X = \begin{cases} \text{"red"} & (9/30) \\ \text{"white"} & (8/30) \\ \text{"green"} & (13/30) \end{cases}$$

$$Y = \begin{cases} \text{"left"} & (10/30) \\ \text{"right"} & (20/30) \end{cases}$$



$$(X|Y=\text{"left"}) = \begin{cases} \text{"red"} & (2/10) \\ \text{"white"} & (1/10) \\ \text{"green"} & (7/10) \end{cases}$$

$$(X|Y=\text{"right"}) = \begin{cases} \text{"red"} & (7/20) \\ \text{"white"} & (7/20) \\ \text{"green"} & (6/20) \end{cases}$$


$$H(X|Y = \text{right}) = H\left(\frac{7}{20}, \frac{7}{20}, \frac{6}{20}\right)$$

$$= -\frac{7}{20} \log_2\left(\frac{7}{20}\right) - \frac{7}{20} \log_2\left(\frac{7}{20}\right) - \frac{6}{20} \log_2\left(\frac{6}{20}\right) = 1.5813$$

Entropy

$$X = \begin{cases} \text{"red"} & (9/30) \\ \text{"white"} & (8/30) \\ \text{"green"} & (13/30) \end{cases}$$

$$Y = \begin{cases} \text{"left"} & (10/30) \\ \text{"right"} & (20/30) \end{cases}$$



$$(X|Y=\text{"left"}) = \begin{cases} \text{"red"} & (2/10) \\ \text{"white"} & (1/10) \\ \text{"green"} & (7/10) \end{cases}$$

$$(X|Y=\text{"right"}) = \begin{cases} \text{"red"} & (7/20) \\ \text{"white"} & (7/20) \\ \text{"green"} & (6/20) \end{cases}$$


$$H(X|Y) = P(Y = \text{left}) \cdot H(X|Y = \text{left}) + P(Y = \text{right}) \cdot H(X|Y = \text{right}) =$$

$$= \frac{10}{30} \cdot 1.1568 + \frac{20}{30} \cdot 1.5813 = 1.4398$$

Information Theory

$$X = \begin{cases} \text{"red"} & (9/30) \\ \text{"white"} & (8/30) \\ \text{"green"} & (13/30) \end{cases}$$

$$Y = \begin{cases} \text{"left"} & (10/30) \\ \text{"right"} & (20/30) \end{cases}$$




$$(X|Y=\text{"left"}) = \begin{cases} \text{"red"} & (2/10) \\ \text{"white"} & (1/10) \\ \text{"green"} & (7/10) \end{cases}$$

$$(X|Y=\text{"right"}) = \begin{cases} \text{"red"} & (7/20) \\ \text{"white"} & (7/20) \\ \text{"green"} & (6/20) \end{cases}$$

$$I(X|Y) = H(X) - H(X|Y) = 1.5524 - 1.4398 = 0.1126 \text{ bit}$$

Information gain - example



$$H(X) = H\left(\frac{9}{18}, \frac{9}{18}\right) = 1$$

$$H(X|Y = \text{left}) = H\left(\frac{8}{12}, \frac{4}{12}\right) = 0.918$$

$$H(X|Y = \text{right}) = H\left(\frac{5}{6}, \frac{1}{6}\right) = 0.650$$

$$H(X|Y) = \frac{12}{18} H(X|Y = \text{left}) + \frac{6}{18} H(X|Y = \text{right}) = 0.8287$$


$$I(X|Y) = H(X) - H(X|Y) = 1 - 0.8287 = 0.1713$$

Most common stopping criterias


- maximum depth
- maximum leaf node
- min samples leaf (only split a node if both children resulting have at least this many samples)
- min samples split (only split a node if it has at least this many samples)

classification example


max_depth = 1



max_depth = 2




max_depth = 3



Calculate the feature importances as follows:

```
feat_imp = [0 for f in features]
for branch in tree.branches:
    feature = branch.feature
    gain = branch.information_gain
    weight = branch.samples_ratio
    feat_imp[f] += weight * gain
feat_imp = normalize(feat_imp)
```

feature importances - example



$f_i = [0, 0, 0]$

$f_i[X_1] += \text{gain} * \text{weight} = 0.1 * 1 = 0.1$
 $f_i = [0.1, 0, 0]$

$f_i[X_1] += 0.1 * 0.6 = 0.06$
 $f_i = [0.16, 0, 0]$

$f_i[X_2] += 0.3 * 0.4 = 0.12$
 $f_i = [0.16, 0.12, 0]$

$f_i[X_3] += 0.2 * 0.5 = 0.1$
 $f_i = [0.16, 0.12, 0.1]$

$f_i[X_1] += 0.1 * 0.48 = 0.048$
 $f_i = [0.208, 0.12, 0.1]$

$f_i = [0.486, 0.28, 0.234]$

pros:

- no need to scale input features
- highly interpretable
- provides feature importance

cons:

- not generalize well
- not so good at picking linear relationships between features
- easy to overfit
- can be computationally expensive
- decision boundaries are parallel to the axes



Converting weak learners into one strong learner.

Boosting algorithms:

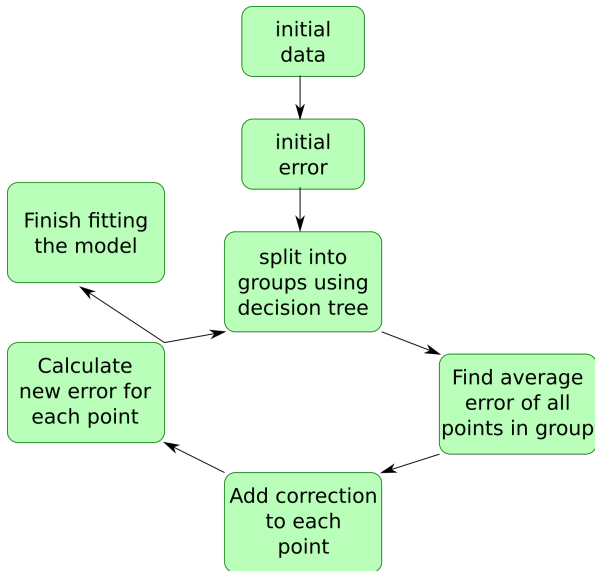
- AdaBoost
- Gradient Tree Boosting
- BrownBoost
- CoBoosting
- GentleBoost
- etc.

why boosting?

It is very powerful. Some competitions that have been won (in Kaggle) using boosting, or boosting in conjunction with other techniques:

- Liberty Mutual Property Inspection
- Caterpillar Tube Pricing
- Avito Duplicate Ads Detection
- Facebook Robot Detection
- Otto Product Classification


Gradient Boosted Trees Regression



Gradient Boosted Trees Regression

- initial prediction can be 0 (or average)
- use weak learners (e.g. decision tree with $\text{max_depth} = 2$)
- $\text{correction} = \text{learning_rate} * \text{error}$


initial data




Apply the boosting algorithm with parameters:

- `learning_rate = 0.5`
- `max_depth = 1`


initial error and groups




iteration 1 - group 1 - corrections




iteration 1 - group 2 - corrections




iteration 2 - current error




iteration 2 - group 1




iteration 2 - group 2




iteration 3







iteration 3 - group 1




iteration 3 - group 2



test errors



test errors



pros and cons

pros:

- powerful at prediction
- harder to overfit

cons:

- slower training (sequential train)
- hard to interpret the predicted value

sklearn.ensemble.GradientBoostingClassifier
sklearn.ensemble.GradientBoostingRegressor
sklearn.ensemble.AdaBoostClassifier
sklearn.ensemble.AdaBoostRegressor
xgboost: Scalable and Flexible Gradient Boosting